

Hojoon Leo Kim

hojoon.kim@snu.ac.kr | hojoonleokim.com

RESEARCH INTERESTS

Machine Learning Systems, Hardware–Software Co-Design, Energy-Efficient AI

My research focuses on agent–hardware co-design: jointly optimizing agentic algorithms for planning and memory with accelerator architectures and execution mechanisms such as plan reuse. I extend this full-stack approach to embodied AI and robotics, where on-device vision–language–action (VLA) models require real-time, energy-efficient inference. My long-term goal is to develop reusable system abstractions and open benchmarks that bridge ML algorithms and accelerator architecture, enabling scalable deployment of agentic and embodied AI.

EDUCATION

• Seoul National University

M.S. in Computer Science and Engineering (Advisor: Professor Jae W. Lee)

B.S. in Electrical and Computer Engineering, Summa Cum Laude

◦ B.S. GPA: 3.99/4.30 (Major GPA: 4.10/4.30)

Seoul, South Korea

Sep 2026 – Aug 2028 (Expected)

Mar 2020 – Feb 2026

• Stanford University

Visiting Student Researcher

Jun 2025 – Aug 2025

Stanford, CA, USA

PUBLICATIONS

C=CONFERENCE, S=IN SUBMISSION

* Equal contribution

- [C.3] Hojoon Kim, Yuheng Wu, and Thierry Tamba. **AgenticCache: Cache-Driven Asynchronous Planning for Embodied AI Agents**. *MLSys'26*. (Acceptance Rate: 26.4%) [Paper PDF]
- [C.2] Seung Yul Lee, Hojoon Kim, Yutack Park, Dawoon Jeong, Seungwu Han, Yeonhong Park, and Jae W. Lee. **FlashTP: Fused, Sparsity-Aware Tensor Product for Machine Learning Interatomic Potentials**. *ICML'25 Spotlight* (Spotlight Rate: 2.6%) [Paper PDF]
- [C.1] Yeonhong Park*, Jake Hyun*, Hojoon Kim, and Jae W. Lee. **DecDEC: A Systems Approach to Advancing Low-Bit LLM Quantization**. *OSDI'25*. (Acceptance Rate: 15.9%) [Paper PDF]
- [S.4] Hojoon Kim, Chanyeong Lim, Jae W. Lee, and Kiwan Maeng. **Carnevale: Hardware-Software Co-Design for Efficient Masked Machine Learning at Scale**. *In submission to ASPLOS'27*
- [S.3] Jaehoon Yang, Yongbeom Kim, Hojoon Kim, Seung Yul Lee, and Jae W. Lee. **FLIT: A Fine-Grained Memory Management Layer for Colocating LLM Inference and Tuning**. *In submission to EuroSys'27*
- [S.2] Jongheon Jeong*, Hojoon Kim*, Rokhee Lee, Yeonhong Park, Young H. Oh, and Jae W. Lee. **Helios: Heterogeneous lightweight VLA model serving system**. *In submission to ASPLOS'27*
- [S.1] Seong Hoon Seo, Donghyun Lee, Geonha Lee, Hojoon Kim, Yeonhong Park, and Jae W. Lee. **QUESO: Storage-Assisted Quantization Error Compensation for On-Device LLM Inference**. *In submission to DATE'27*

RESEARCH EXPERIENCE

• SNU Architecture and Code Optimization Lab [ARC Lab]

Graduate Researcher (P.I.: Professor Jae W. Lee)

Undergraduate Researcher

Seoul, South Korea

Feb 2026 – Aug 2028 (Expected)

Sep 2024 – Feb 2026

- Led the compiler–architecture co-design for exact INT64 arithmetic, developing optimized GPU kernels and proposing GPU microarchitectural extensions for scalable masked machine learning. (ASPLOS'27 submission)
- Led the design and implementation of HELIOS, a lightweight GPU kernel orchestrator for centralized multi-robot VLA serving. (ASPLOS'27 submission)

• Stanford Tambe Lab [Tambe Lab]

Undergraduate Researcher (P.I.: Professor Thierry Tambe)

- Led research on cache-driven asynchronous planning for embodied AI agents by exploiting plan-level locality. (MLSys'26)

Jun 2025 – Mar 2026

Stanford, CA

• SNU Computer Architecture and Systems Lab [COMPARCH Lab]

Research Intern (P.I.: Professor Jaewoong Sim)

- Authored undergraduate thesis based on [Tender \(ISCA'24\)](#) quantization architecture; designed hardware module to shift offline computation to runtime, improving model accuracy.

Dec 2023 – Sep 2024

Seoul, South Korea

ACADEMIC PROJECTS

- **A File System with Full-Path Indexing: Operating Systems (1st Place)** Dec 2024
Reduced file lookup latency from 24,000 to 993 cycles on the xv6-riscv kernel
- **Pipelined Central Processing Unit with Perceptron Branch Predictor: Computer Organization** Jun 2024
- **L1, L2, and Memory Three-Level Memory Simulator with MSHR: Computer Organization** Jun 2024
- **A CNN Accelerator on FPGA (Artix-7): Digital Systems Design and Lab (1st Place)** Dec 2023

HONORS AND AWARDS

- **YoulChon AI Young Researcher Scholarship** Sep 2026
Yulchon Foundation & Seoul National University AI Institute
- **Semiconductor Specialized Scholarship** Jan 2024 – Feb 2026
Semiconductor Specialization Project Group, Seoul National University
- **Grand Prize – 2025 AI Chip Contest - NPU Optimization Track** Nov 2025
Ministry of Science and ICT (MSIT), Republic of Korea
- **Global Leadership Program (airfare)** Jun 2025 – Aug 2025
College of Engineering
- **Merit-Based Tuition Scholarship** Aug 2023 – Jun 2025
Department of Electrical and Computer Engineering
- **1st Prize – Advanced CUDA Accelerator Programming School** Feb 2025
SNU Thunder Research Group
- **4th Prize – Military AI Contest** Dec 2022
Ministry of National Defense & Ministry of Science and ICT

ACADEMIC SERVICE

- **Program Committee Member, MLSys 2027** 2027
Conference on Machine Learning and Systems

TEACHING EXPERIENCE

- **Peer Tutor, Computer Organization; Programming Methodology** Sep 2024 – Jun 2025
Department of Electrical and Computer Engineering

MILITARY SERVICE

- **SIGINT Specialist, Republic of Korea Army** Jan 2022 – Jul 2023
 - Two-time recipient of the Best SIGINT Analyst Award; recognized as a First-Class Soldier (fitness, mental readiness, and specialty proficiency).

EXTRACURRICULAR ACTIVITIES

- **Member, SNU Business School Tennis Club** Mar 2024 – Dec 2024
- **Peer Mentor, SNU Center for Campus Life & Culture** Mar 2024 – Jun 2024
 - Received the Best Mentor Award and Best Mentoring Team Award.

TECHNICAL SKILLS

- **Programming Languages:** C/C++, Python, Verilog/SystemVerilog
- **ML Systems and GPU Programming:** CUDA, PyTorch, NVIDIA Nsight Systems/Compute
- **EDA and FPGA Tools:** Synopsys Design Compiler (logic synthesis, area estimation), Synopsys PrimeTime (timing/power analysis), Xilinx Vivado
- **Architecture Simulators:** Accel-Sim, GPGPU-Sim

LANGUAGE PROFICIENCY

Korean (Native), English (TOEFL 106 [R:30/L:28/S:22/W:26], OPIc AL)