

Hojoon Leo Kim

hojoon.kim@snu.ac.kr | hojoonleokim.com

RESEARCH INTERESTS

Machine Learning Systems, Efficiency, HW-SW Co-Design

As a researcher passionate about advancing ML systems through hardware-software co-design, I explore how traditional computer architecture principles such as caching, branch prediction, and multiprocessing can address inefficiencies in modern ML workloads. My work spans low-bit quantization, storage-assisted inference, and AI agents, with publications at OSDI'25, ICML'25 (Spotlight), and MLSys'26. I aim to develop practical system architectures that make next-generation ML applications both efficient and deployable at scale by rethinking system abstractions across the entire computing stack.

EDUCATION

- **Seoul National University** Mar 2020 – Feb 2026
Undergraduate, Electrical and Computer Engineering Seoul, South Korea
 - *Summa Cum Laude*
 - GPA: 3.99/4.30 (major GPA: 4.10/4.30)
- **Stanford University** Jun 2025 – Aug 2025
Visiting student researcher Stanford, CA, USA
 - GPA: 4.15/4.30; Courses: CS107 - Computer Organization and Systems (A+), CS161 - Algorithm (A0)

PUBLICATIONS

C=CONFERENCE, S=IN SUBMISSION

* Equal contribution

- [S.2] Jongheon Jeong*, Hojoon Kim*, Rokhee Lee, Yeonhong Park, Young H. Oh, and Jae W. Lee. **Helios: Heterogeneous lightweight VLA model serving system**. *In submission to EuroSys'27*
- [S.1] Seong Hoon Seo, Donghyun Lee, Geonha Lee, Hojoon Kim, Yeonhong Park, and Jae W. Lee. **QUESO: Storage-Assisted Quantization Error Compensation for On-Device LLM Inference**. *In submission to ATC'26*
- [C.3] Hojoon Kim, Yuheng Wu, and Thierry Tambe. **AgenticCache: Cache-Driven Asynchronous Planning for Embodied AI Agents**. *MLSys'26*. (Acceptance Rate: 26.4%) [\[Paper PDF\]](#)
- [C.2] Seung Yul Lee, Hojoon Kim, Yutack Park, Dawoon Jeong, Seungwu Han, Yeonhong Park, and Jae W. Lee. **FlashTP: Fused, Sparsity-Aware Tensor Product for Machine Learning Interatomic Potentials**. *ICML'25 Spotlight* (Spotlight Rate: 2.6%) [\[Paper PDF\]](#)
- [C.1] Yeonhong Park*, Jake Hyun*, Hojoon Kim, and Jae W. Lee. **DecDEC: A Systems Approach to Advancing Low-Bit LLM Quantization**. *OSDI'25*. (Acceptance Rate: 15.9%) [\[Paper PDF\]](#)

RESEARCH AND WORK EXPERIENCE

- **SNU Architecture and Code Optimization Lab [\[ARC Lab\]](#)** Seoul, South Korea
Post-Bachelor Researcher (P.I.: Professor Jae W. Lee) Feb 2026 – Present
Undergraduate Researcher Sep 2024 – Feb 2026
 - Researched error compensation at prefill stage with low-precision decoding to minimize GPU memory usage while improving model accuracy. (MLSys'26 submission)
 - Conducted figure design, experimental evaluation, and kernel performance analysis for Tensor-Product acceleration project. (ICML'25 Spotlight)
 - Built AWQ-quantized 3/3.5/4-bit LLMs using lutgemm-based weight packing (bitplane-wise packing + GEMV) for dynamic error compensation. (OSDI'25)
- **Stanford Tambe Lab [\[Tambe Lab\]](#)** Jun 2025 – Mar 2026
Undergraduate Researcher (P.I.: Professor Thierry Tambe) Stanford, CA
 - Conducted research on cache-driven asynchronous planning for embodied AI agents by exploiting plan-level locality. (MLSys'26 submission)
- **SNU Computer Architecture and Systems Lab [\[COMPARCH Lab\]](#)** Dec 2023 – Sep 2024
Research Intern (P.I.: Professor Jaewoong Sim) Seoul, South Korea
 - Authored undergraduate thesis based on [Tender \(ISCA'24\)](#) quantization architecture; designed hardware module to shift offline computation to runtime, improving model accuracy.
 - Applied delayed-aggregation (Mesorasi) to point cloud models to analyze throughput and accuracy.
- **MODULABS [\[MODULABS\]](#)** Jun 2021 – Sep 2021
Backend Software Engineer Seoul, South Korea
 - Developed backend for AI education platform: implemented payment processing system with time-synchronized functions and managed backend database infrastructure for AIFEL, a new business initiative.

ACADEMIC PROJECTS

- **A File System with Full-Path Indexing: Operating Systems** Dec 2024
 - Achieved 1st place (evaluated by throughput). Extended the [xv6-riscv OS kernel](#) to support full-path indexing and faster file access. Optimizations included (1) hash-based file lookup and (2) a dedicated sleep-wait stack to reduce scheduling overhead. Reduced file lookup latency from 24,000 to 993 cycles.
- **Pipelined Central Processing Unit with Perceptron Branch Predictor: Computer Organization** Jun 2024
- **L1, L2, and Memory Three-Level Memory Simulator with MSHR: Computer Organization** Jun 2024
- **A CNN Accelerator on FPGA(ARTIX-7): Digital Systems Design and Lab** Dec 2023
 - Ranked 1st based on throughput and task completion time. Optimized performance by addressing memory and compute bottlenecks: (1) implemented multi-bank BRAM architecture and DMA-driven dataflow to fetch data in parallel, resolving memory-bound issues; (2) used multiple output-stationary systolic arrays to boost compute throughput; and (3) pipelined the Load and Compute states to enable weight prefetch and overlap memory transfer with computation.

HONORS AND AWARDS

- **SNU Semiconductor Specialized Scholarship (KRW 15,000,000)** Jan 2024 – Feb 2026
SNU Semiconductor Specialization Project Group
- **Grand Prize (1st Place), 2025 AI Chip Contest - NPU Optimization Track (KRW 10,000,000)** Nov 2025
Ministry of Science and ICT (MSIT), Republic of Korea
- **Global Leadership Program (airfare)** Jun 2025 – Aug 2025
College of Engineering
- **Merit-Based Tuition Scholarship (70% of tuition)** Aug 2023 – Jun 2025
Department of Electrical and Computer Engineering
- **1st Prize – Advanced CUDA Accelerator Programming School** Feb 2025
SNU Thunder Research Group
- **4th Prize – Military AI Contest (KRW 2,000,000)** Dec 2022
Ministry of National Defense & Ministry of Science and ICT

TEACHING EXPERIENCE

- **Peer Tutor, Computer Organization; Programming Methodology** Sep 2024 – Jun 2025
Department of Electrical and Computer Engineering

EXTRACURRICULAR ACTIVITIES

- **Member, SNU Business School Tennis Club** Mar 2024 – Dec 2024
- **Peer Mentor, SNU Center for Campus Life & Culture** Mar 2024 – Jun 2024
 - Received the Best Mentor Award and Best Mentoring Team Award.
- **SIGINT Specialist, Republic of Korea Army** Jan 2022 – Jul 2023
 - Two-time recipient of the Best SIGINT Analyst Award; recognized as a First-Class Soldier (fitness, mental readiness, and specialty proficiency).

RELEVANT COURSES

- | | |
|--|--|
| • System Programming (Prof. Jae W. Lee) | • Advanced Compiler (Prof. SooMook Moon) |
| • Operating Systems (Prof. JinSoo Kim) | • Computer Networks (Prof. KyoungSoo Park) |
| • Computer Organization (Prof. Jaewoong Sim) | • Linear Algebra (Prof. Young-Hoon Kiem) |
| • Advanced Computer Architecture (Prof. Sungjoo Yoo) | • Parallel Distributed Computing (Prof. Jinho Lee) |
| • Digital Systems Design (Prof. Jaewoong Sim) | • Computer Organization & Design (Prof. Jangwoo Kim) |

TECHNICAL STRENGTHS

- **Computer Languages:** C/C++, Verilog/SystemVerilog, Python
- **Frameworks and Skills:** CUDA, PyTorch, Nsight System/Compute

LANGUAGE PROFICIENCY

Korean (Native), English (TOEFL 106 [R:30/L:28/S:22/W:26], OPIc AL)